

Alignment between CLS attention and Gradient×Input saliency across layers
(mean $\pm$ std over 399 correctly classified test images)

